

Early co-occurrence network openness predicts cross-disciplinary concept breadth: evidence from 12,499 scientific concepts

AI Inventor

Abstract

Which early structural signals in a concept’s topic co-occurrence network anticipate later cross-field breadth? Using the OpenAlex bulk snapshot, we identify 12,499 concepts across 26 scientific fields and screen 53 early network indicators against a five-feature popularity baseline on held-out field groups. Seven indicators survive held-out validation and compose into an openness index (OPEN) that measures a concept’s early tendency to acquire novel, community-spanning co-occurrence partners. A descriptive meta-analysis of $\text{OPEN}_{\text{home}}$ over six non-selection bodies gives partial Spearman $+0.069$ [$+0.038, +0.100$], $I^2 = 0$, all six positive; the evidence state is lead. The association behaves like a static topical-dispersion property: a within-concept permutation null absorbs year-to-year churn, and the temporal-excess measure has split-half reliability below 0.05, so temporal turnover is untested rather than refuted. On an independent frame of 11,841 concepts, a conditional logit on field-entry events shows that concepts spread next to fields related to the ones currently retaining them ($d_0 = 0.322$ [$0.291, 0.355$]), but the pre-registered verdict is PARTIAL: the volume-matched contrast is null and the coefficient reverses under an alternative proximity measure. Early contact diversity accounts for most of the breadth gap; no discrete trajectory typology passes the naming rule.

Keywords: co-occurrence networks, cross-disciplinary diffusion, knowledge spread, principle of relatedness, scientific concepts, temporal networks, topic emergence

1 Introduction

Some scientific concepts remain confined to their home discipline for decades, while others cross field boundaries within a few years. Optogenetics, born in neuroscience, entered genetics, psychiatry and bioengineering; deep learning, rooted in computer science, now appears in medicine, materials science and linguistics. Understanding what distinguishes broadly spreading concepts from locally absorbed ones matters for science policy, research evaluation and the design of interdisciplinary programmes [Fortunato et al., 2018, Hofstra et al., 2019]. If early structural signals predict later cross-field integration, funders and institutions could identify concepts with high diffusion potential before that diffusion is observable in publication counts alone.

The problem is hard because simple popularity metrics (publication volume, growth rate, share of papers outside the home field) already explain much of the variance in later breadth [Rotolo et al., 2015]. Any network indicator must add information *beyond* these baselines. Network science offers a rich toolkit for studying how information, behaviour and innovations spread through structured populations [O’Brien et al., 2018]. In citation and co-authorship networks, knowledge diffusion follows structured pathways that reflect both the topology of the network and the position of the spreading entity within it [Domenico et al., 2016, Renoust et al., 2017]. Community structure shapes diffusion outcomes: bridge nodes and boundary-spanning ties accelerate spread across

groups [Cherifi et al., 2019, Burt, 2004], while community cohesion can either facilitate or retard it depending on the contagion mechanism [Centola, 2010, Ugander et al., 2012].

Prior work has proposed several candidate signals for concept diffusion. Cheng et al. [2023] tracked roughly 2,000 new ideas across Web of Science fields and found that early social reach, the number of disconnected researcher groups, was the strongest predictor of whether an idea became “core.” Their ideational consistency measure predicted next-year volume. Weng et al. [2013] showed that early community diversity predicts virality of online content (AUC 0.83). Salatino et al. [2017] showed that collaboration density rises before topic emergence, and Chavalarias and Cointet [2013] used phylomemetic methods to track the branching and merging of scientific fields. Gao et al. [2018] tracked community evolution in patent co-occurrence networks, revealing how technological change reshapes network structure. Sattar et al. [2023] developed approaches for detecting the temporal evolution of communities, and Callon et al. [1991] introduced co-word analysis as a tool for mapping research structure. These studies typically operate on curated concept sets without held-out validation, mix concept-level and field-level units of analysis, or do not control for simple volume-based predictors. No prior study has screened a broad set of temporal knowledge-network indicators against a shared baseline on held-out field groups for the specific outcome of size-adjusted cross-disciplinary breadth.

A second open question concerns the routes by which concepts move between fields. The principle of relatedness, established for regional economic diversification [Hidalgo et al., 2007, Neffke et al., 2011, Fitzgerald et al., 2021], predicts that actors diversify into activities related to their existing portfolio. For scientific fields adopting new concepts, this principle has been tested for entry but not for the role of currently *retaining* fields—those that have adopted a concept and kept publishing on it. Whether concepts spread next to fields related to the ones currently retaining them, beyond relatedness to the home field alone, is untested.

This paper addresses both questions with a large-scale empirical study. We identify 12,499 concepts from the OpenAlex bulk snapshot, compute 53 early network indicators across six families, and test them on held-out field groups and a confirmatory onset cohort. We also build a conditional logit model of field entry that tests whether retained-field relatedness predicts the next field a concept enters. We report the following contributions:

1. We screen 53 early co-occurrence network indicators¹ against a five-feature popularity baseline and confirm seven on held-out field groups for predicting rarefied cross-field breadth. Two of the seven carry negative signs: early retention ratio and ego-network density.
2. We introduce the OPEN index, a six-component composite of early co-occurrence openness, and test it on a 2015–2017 cohort never used in selection. A descriptive meta-analysis of $OPEN_{home}$ over six non-selection bodies gives partial Spearman $+0.069$ [$+0.038, +0.100$], $I^2 = 0$, 6/6 positive; the evidence state is lead.
3. We test whether concepts spread next to fields related to the ones currently retaining them. On an independent frame² (11,841 concepts, 6,978 entry events in 6,076 informative strata), the retained-frontier coefficient is $d_0 = +0.322$ [$+0.291, +0.355$]. The pre-registered verdict is FRONTIER = PARTIAL: the volume-matched contrast is null, the coefficient reverses under minimum-conditional-probability proximity, and the held-out dose response is not monotone. Persistence and volume remain confounded.

¹Code: https://github.com/ai-inventor-papers/ai-invention-8795ea-concepts-spread-once-adopters-make-them/tree/fork/run_DVtwCx0JbFq/round-3/experiment-8

²Code: https://github.com/ai-inventor-papers/ai-invention-8795ea-concepts-spread-once-adopters-make-them/tree/fork/run_DVtwCx0JbFq/round-3/experiment-7

4. We show that early neighbourhood consolidation, measured by Cheng et al.’s (2023) ideational consistency, predicts volume growth but *narrower* cross-field reach, a reversal that holds in all five field groups.
5. We decompose breadth into exploration and retention channels and show that early contact diversity accounts for 73% of the breadth gap. No discrete trajectory typology passes the naming rule; diffusion paths form a continuum along a breadth axis.

2 Methods

2.1 Concept identification and panel construction

We use the OpenAlex bulk snapshot (2026-09-23; 476,196,327 works; 129.4 million base works 1995–2022). Concepts are identified by Aho–Corasick title matching of 56,643 legacy OpenAlex concepts (levels 2–5) plus Wikidata aliases, with stemmed verification. A per-concept LLM precision gate drops concepts with precision below 0.80.

A concept’s *onset year* t_0 is the first year it reaches 20 grounded publications. *Home field(s)* are the field(s) holding at least 40% of the first 30 grounded papers. The specification was frozen on DEV data (hash-sealed before held-out scoring) and unsealed once. The panel comprises 12,499 concepts and 27,393 concept-by-field adoption episodes (Table 1).

Table 1: Panel composition. DEV concepts have home fields in Computer Science, Engineering, Biochemistry/Genetics/Medicine. Held-out groups cover the remaining field families.

Split	Concepts	Episodes
DEV (CS/Eng/BGM/Med, onset 2003–2009)	4,771	9,079
Cohort (onset 2010–2014, all fields)	4,356	9,799
Held-out Physical Sciences	742	1,662
Held-out Life & Environment	1,113	3,099
Held-out Social Sciences	1,352	3,320
Held-out Math & Decision Sci.	165	434
Total	12,499	27,393

2.2 Outcome: rarefied field breadth

The primary outcome is *rarefied field breadth* O_{2r} ($m = 50$): the expected number of distinct venue fields among a fixed-size random draw of m papers from a concept’s publications in years $t_0 + 6$ to $t_0 + 8$, computed by exact hypergeometric rarefaction. Rarefaction separates size-adjusted breadth from sheer volume. Secondary outcomes include sustained uptake (O_{1c}), transience (O_3), and field- and year-normalised citation growth (O_4).

2.3 Baseline and indicator families

We define a *five-feature baseline* (B5) consisting of log early volume, publication growth, non-home share, Shannon entropy and field reach, all computed over t_0 to $t_0 + 2$. The 53 candidate indicators span six families:

1. **Co-occurrence ego-network (A, 27 indicators):** degree, density, persistence, novelty, community counts, participation, brokerage from the concept’s ego network on the co-occurrence backbone.
2. **Popularity and volume (E, 6 indicators):** publication growth, author growth, author count.
3. **Disciplinary spread (F, 3 indicators):** Shannon entropy of field distribution, home-field relatedness, Rao–Stirling diversity.
4. **Retained frontier (FR, 7 indicators):** contact reach, retention ratio, RCA-based and volume-weighted field densities.
5. **Gateway centrality (G, 7 indicators):** eigenvector centrality of home and early off-home fields on the field backbone.
6. **Co-author reach (S, 3 indicators):** author overlap with other concepts, connected components among co-authors.

2.4 The OPEN index

The OPEN index is the mean of six signed z -scored ego-network components from the early window (t_0 to $t_0 + 2$):

1. New edge rate (positive sign): the fraction of a concept’s co-occurrence neighbours that are new.
2. Number of Leiden communities reached (n_{comm} , positive): how many distinct topic communities a concept’s neighbours span.
3. Participation coefficient (positive): the share of a concept’s ties that cross community boundaries.
4. Neighbourhood novelty residual (NOV_{res} , positive): the novelty of a concept’s co-occurrence partners after residualising on degree.
5. Ego density (negative sign): the density of the concept’s local co-occurrence subgraph.
6. Edge persistence (negative sign): the fraction of the concept’s neighbours retained across windows.

The z -score constants are frozen from the DEV data. We test three builds: $\text{OPEN}_{\text{home}}$ (home-field co-occurrence graph only), OPEN_{all} (corpus-wide graph), and $\text{OPEN}_{\text{sizematch}}$ (corpus-wide papers subsampled to home counts).

2.5 Statistical framework

Indicator screening uses *partial Spearman priority* (PSP): the Pearson correlation of rank-transformed indicator and outcome residuals, both residualised on the baseline covariates and control rungs. Inference uses 2,000 concept-level bootstrap resamples with refit. Pooling across field groups uses the DerSimonian–Laird random-effects estimator [DerSimonian and Laird, 1986] with I^2 heterogeneity reporting [Higgins and Thompson, 2002]. Multiple comparisons use Holm correction across the pre-declared family.

We also define NOVCHURN as the mean of $z(\text{NOV}_{\text{res}})$ and $-z(\text{edge persistence})$, a two-component novelty/churn measure used in mechanism analyses.

2.6 Field-entry model (RQ2)

For the field-entry analysis, we use a conditional logit (Breslow partial likelihood) on concept-year risk sets, stratified by concept-year, with standardised covariates: relatedness to home, log field size, density over entered fields, gateway centrality (baseline R0); RCA-based density (R1); volume-weighted density (R2); and retained-field relatedness d_0 (R3). A field is “retaining” a concept when it has published at least two grounded papers on the concept in each of the two most recent years. Relatedness is measured by positive pointwise mutual information (PMI) on the field-level backbone. The specification is frozen on DEV and scored once on the independent held-out frame.

3 Results

The headline finding is the openness–breadth association: a descriptive meta-analysis of $\text{OPEN}_{\text{home}}$ over six non-selection bodies gives partial Spearman $+0.069$ [$+0.038, +0.100$], $I^2 = 0$, 6/6 positive; the evidence state is lead. We present the indicator screen (RQ1) first because it defines the OPEN index.

3.1 RQ1: Which early network signals predict cross-field breadth?

3.1.1 Experimental setup

The top 10 indicators selected on DEV by partial Spearman priority for rarefied breadth (O_{2r} , $m = 50$) were tested once on four held-out field groups and two cohort splits. Confirmation requires Holm-corrected $p < 0.05$ and a 95% CI excluding zero on the DerSimonian–Laird pooled estimate. The held-out outcomes had been previously used by two earlier experiments, so the indicator-screen results are robustness checks within the same frame rather than fully independent confirmations. The fresh 2015–2017 cohort is the only body never used in any prior step.

3.1.2 Comparison to related work

Prior network-based diffusion studies in *Applied Network Science* have examined knowledge flow through citation and co-authorship networks [Domenico et al., 2016, Renoust et al., 2017], community-driven diffusion mechanisms [Cherifi et al., 2019, O’Brien et al., 2018], and temporal community evolution [Sattar et al., 2023, Gao et al., 2018]. These studies model diffusion at the network level but do not screen concept-level co-occurrence indicators against a held-out validation framework for the specific outcome of size-adjusted breadth. Weng et al. [2013] showed that early community diversity predicts online virality (AUC 0.83); our study extends this to scientific concepts, where the baseline is much stronger (B5 Spearman 0.65–0.86). Cheng et al. [2023] found that social reach predicts whether ideas become “core,” but did not control for a shared baseline or use held-out field groups. Fitzgerald et al. [2021] showed that academic knowledge is becoming more localised, motivating the question of which early signals distinguish concepts that overcome this localisation.

3.1.3 Results

Seven of the 53 screened indicators survive held-out validation with Holm correction (Table 2). These span two families: retained frontier (end-of-window field density, volume-weighted field den-

sity, contact reach, retention ratio) and co-occurrence topology (community count, neighbourhood novelty residual, ego density). No centrality or volume-only indicator survives. Two confirmed indicators carry negative signs: the early retention ratio (-0.114 , concepts with higher early field retention spread *less* broadly) and ego density (-0.102 , concepts with denser co-occurrence neighbourhoods spread less).

Table 2: Held-out indicator screen: the 10 indicators with highest DEV partial Spearman priority for rarefied breadth (O_{2r} , $m = 50$), tested on held-out groups. DerSimonian–Laird pooled estimates with 95% CI. Confirmed indicators (Holm $p < 0.05$ and CI excluding zero) are marked.

Indicator	Family	Pooled PSP	95% CI	I^2	Sign	Conf.
M0_density_end	FR	+0.375	[+0.279, +0.462]	0.74	6/6	Yes
D_vol_end	FR	+0.307	[+0.256, +0.356]	0.10	6/6	Yes
CONTACT_REACH	FR	+0.211	[+0.161, +0.261]	0.00	6/6	Yes
n_{comm} (W3)	A	+0.167	[+0.063, +0.267]	0.78	6/6	Yes
NOV _{res}	A	+0.151	[+0.044, +0.255]	0.75	6/6	Yes
RETENTION_RATIO	FR	-0.114	[-0.160 , -0.067]	0.00	6/6	Yes
ego_density (W3)	A	-0.102	[-0.151 , -0.053]	0.00	6/6	Yes
Rao–Stirling	G	-0.072	[-0.153 , +0.010]	0.44	5/6	No
G_{btw}	G	+0.056	[-0.006 , +0.118]	0.33	6/6	No
log offhome vol.	F	-0.089	[-0.171 , -0.007]	0.63	5/6	No

About half of the M0_density_end and D_vol_end signal is a pre-onset footprint: on held-out data, the post-onset-only field density attenuates from +0.374 to +0.187 [+0.145, +0.246]. The remaining indicators measure genuinely early structural properties.

An ElasticNet combining all indicators adds Spearman +0.059 [+0.046, +0.073] over B5 on held-out groups; an Explainable Boosting Machine adds +0.052 [+0.037, +0.067].

OPEN index confirmation. The OPEN index was tested on a 2015–2017 onset cohort³ (1,443 concepts) never used in any prior screen. We define a six-rung control ladder that tests OPEN’s partial Spearman with rarefied breadth under increasingly demanding baselines (Table 3).

Table 3: OPEN index control ladder on the 2015–2017 confirmatory cohort. PSP = partial Spearman priority with rarefied breadth (O_{2r} , $m = 50$), 95% concept-bootstrap CI.

Build	R0	R1	R2	R3	R4	R5	n
OPEN _{home}	+12	+10	+09	+08	+07	+06	573
OPEN _{all}	+21	+18	+17	+17	+15	+14	630
OPEN _{size}	+18	+15	+15	+14	+12	+11	591

OPEN_{home} passes the primary concept-type rung (R2: PSP = +0.091, 95% CI [+0.013, +0.171], Holm $p = 0.048$) but its DerSimonian–Laird pooled interval includes zero (+0.083 [-0.007 , +0.173]), making the home-only signal marginal. The paired difference OPEN_{all} minus OPEN_{home} at the footprint rung is +0.093 [+0.016, +0.169]; about half of this gap is paper count (OPEN_{size} minus OPEN_{home} = +0.053 [-0.015 , +0.117]). OPEN_{all} is partly mechanically coupled to the outcome because early off-home papers that enter the ego network are also papers that determine later breadth.

³Code: https://github.com/ai-inventor-papers/ai-invention-8795ea-concepts-spread-once-adopters-make-them/tree/fork/run_DVtwCx0JbFq/round-4/experiment-10

An evidence synthesis across six non-selection bodies (four held-out domain groups plus the 2010–2014 and 2015–2017 cohorts) gives a DerSimonian–Laird pooled $\text{OPEN}_{\text{home}}$ partial Spearman of $+0.069$ [$+0.038$, $+0.100$] with $I^2 = 0$, all six bodies positive. The DEV estimate ($+0.109$) shows shrinkage of $1.58\times$ relative to the non-selection pool.




Figure 1: Partial Spearman correlation of $\text{OPEN}_{\text{home}}$ (home-neighbourhood openness over t_0 to t_0+2) with rarefied cross-field breadth (O_{2r} , $m = 50$) at ladder rung R2. R2 controls for the five-feature popularity baseline B5 plus onset year, contact reach, concept type, a generic-concept flag and concept level. Blue squares with whiskers show the estimate and 95% concept-bootstrap CI (2,000 draws) for each of six non-selection bodies: four held-out home groups (Physical Sciences, Life & Environment, Social Sciences, Math & Decision; onsets 2003–09), the 2010–14 onset cohort and the 2015–17 onset cohort. n is the number of concepts analysed. The black diamond spans the DerSimonian–Laird pooled estimate over these six bodies, $+0.069$ [$+0.038$, $+0.100$], $I^2 = 0$, 6/6 positive. The dotted vertical line marks this pooled estimate and the dashed line marks zero. The grey row below the separator is the DEV selection body, on which the index was chosen ($+0.109$ [$+0.073$, $+0.144$]). It is not pooled, and it is $1.58\times$ the non-selection pool. Values at right are estimate [95% CI]. The association is small and has the same sign in every body, but only the two cohort rows and the pool have CIs above zero individually. All non-selection bodies except the 2015–17 cohort had outcomes read in earlier analyses, so the pool is descriptive rather than confirmatory.

Vocabulary-free replication (Frame N). To address survivorship bias in the legacy concept vocabulary, we mined 636 vocabulary-free phrase-born concepts⁴ (2003–2015 onsets). These newborns have 14% lower breadth and 89% higher transience than legacy concepts. On Frame N, $\text{OPEN}_{\text{home}}$ gives PSP $+0.117$ [$+0.020$, $+0.218$] at R3 for O_{2r} ($m = 30$); the frozen verdict is PARTIAL: the R5 confidence interval includes zero ($+0.086$ [-0.009 , $+0.190$]), three of four estimable groups are positive, and the Holm-corrected p -value is 0.052. Edge persistence is null on Frame N (-0.013); the signal is carried by neighbourhood novelty ($\text{NOV}_{\text{res}} +0.208$ [$+0.113$, $+0.303$]).

The consistency–breadth reversal. We rebuilt Cheng et al.’s (2023) ideational consistency on OpenAlex topic co-usage for 12,311 concepts. In their design (next-year volume, no current-size

⁴Code: https://github.com/ai-inventor-papers/ai-invention-8795ea-concepts-spread-once-adopters-make-them/tree/fork/run_DVtwWCx0JbFq/round-5/experiment-13

control), the model reproduces their estimate: $b = 0.428$ (+53.5%/SD). Adding current volume removes almost all of it: +1.3% [+0.5%, +2.1%], a residual-to-raw ratio of 0.021. Net of B5, the partial Spearman of consistency with rarefied breadth is -0.069 [-0.093 , -0.047] on the pooled bodies ($I^2 = 0$, 5/5 groups negative), replicating on the 2015–2017 cohort (-0.111 [-0.197 , -0.030]). The consistency measure has Spearman $+0.77$ with edge persistence (Jaccard similarity); it is essentially weighted persistence, confirming that a stable neighbourhood predicts growth but not breadth.

3.1.4 Discussion of RQ1

The confirmed indicators divide into two groups. Contact reach, retention ratio and the field-density measures capture a concept’s position in the inter-field backbone and partly reflect a pre-onset footprint. Community count, neighbourhood novelty and ego density capture the topology of the early home-field co-occurrence neighbourhood. The OPEN index combines these into a single measure of early openness. The practical forecasting gain is small: OPEN_{home} adds only +0.002 Spearman on the fresh cohort. The indicators characterise structural preconditions for broad diffusion rather than providing individually actionable forecasts. Life & Environment Sciences show the weakest OPEN signal (PSP +0.042 vs pooled +0.069), warranting domain-specific investigation.

3.2 RQ2: By what routes do concepts spread across fields?

3.2.1 Experimental setup

We test whether relatedness to the fields currently retaining a concept predicts which field it enters next, using a conditional logit on concept-year risk sets with an independent frame of 11,841 concepts (6,978 entry events in 6,076 informative strata). The frame excludes all concepts used in the original frame construction. The specification was hash-sealed on DEV before held-out scoring. Pre-registered criteria include: pooled R3 CI above zero, positive in the majority of evaluable groups, volume-matched contrast CI above zero, retained-label permutation null rejected, backbone-rewiring null rejected, and dose-response non-negative.

3.2.2 Comparison to related work

Hidalgo et al. [2007] showed that product-space relatedness predicts economic diversification; the minimum conditional probability defines the standard proximity. Neffke et al. [2011] extended the principle of relatedness to regional and technological diversification, and Fitzgerald et al. [2021] showed how regional knowledge networks shape academic localisation. Our contribution is to test whether relatedness to the set of fields currently *retaining* a concept predicts entry, beyond relatedness to the home field and conventional density measures. The key methodological distinction is that retention, not mere presence, defines the reference set.

3.2.3 Results

On the independent held-out frame, the retained-frontier coefficient is $d_0 = 0.322$ [0.291, 0.355] with concept-clustered standard errors (R3 vs R2 likelihood-ratio 325.8, $p < 10^{-70}$). The result is positive in three of three evaluable held-out groups: Physical Sciences +0.148, Life & Environment +0.401, Social Sciences +0.297. Mathematics & Decision Sciences is null (+0.065 [-0.109 , 0.239], 165 concepts). DerSimonian–Laird pooled $d_0 = 0.243$ [0.118, 0.368], $I^2 = 0.92$ (Table 4).

The retained-frontier coefficient survives the conventional RCA density (D_{rca}) and a share-weighted current presence density: adding d_0 after both rivals gives LR = 34.5 on DEV, and D_{rca} is

Table 4: Retained-frontier coefficient (d_0) from the conditional-logit model of field entry, by domain group and cohort. The pre-registered verdict is PARTIAL: criterion 5 (volume-matched CI > 0) fails.

Model	d_0	95% CI (concept)	LR
Development set	0.228	[0.164, 0.291]	34.5
Held-out pooled	0.322	[0.291, 0.355]	325.8
Physical Sciences	0.148	[0.078, 0.219]	–
Life & Environment	0.401	[0.342, 0.460]	–
Social Sciences	0.297	[0.246, 0.348]	–
Math & Decision Sci.	0.065	[−0.109, 0.239]	–
2010–2014 cohort	0.321	[0.292, 0.347]	–

absorbed once d_0 enters. All three specificity tests reject their nulls after Holm correction (retained-label permutation $p = 0.001$; backbone-rewiring $p = 0.004$; node-label $p = 0.003$).

However, the pre-registered verdict is PARTIAL. The volume-matched contrast (comparing entry rates of retained versus entered-but-not-retained fields in the same volume cell) is null: -0.028 $[-0.105, +0.046]$ on held-out data (-0.008 on DEV; Holm $p = 0.76$). The dose response by retention age on held-out data is $0.098 / 0.075 / 0.304$, which is not monotone. Under minimum-conditional-probability proximity [Hidalgo et al., 2007] instead of PMI, $d_0 = -0.021$ ($p = 0.012$), while RCA density becomes strong (LR 246). Persistence is therefore confounded with volume, and the result is backbone-specific.

Gateway centrality. Gateway centrality does not predict field retention: on 27,393 held-out episodes, the increment is effectively zero ($\Delta\text{AUC} -0.00001$, 95% CI $[-0.0006, +0.0003]$).

Breadth decomposition. On 3,188 DEV concepts⁵, the breadth gap between the top and bottom O_{2r} terciles decomposes as: early contact diversity 77.9%, frontier advance -4.6% (broad concepts gain new fields at similar rates), and retention 26.8%. The difference between exploration and retention is 0.464 $[0.407, 0.528]$. On the held-out pool excluding Medicine (variant iv), the contact share is 0.633 $[0.537, 0.727]$ and the difference is 0.504 $[0.329, 0.679]$ ($I^2 = 0.76$). These shares are an accounting identity for the breadth outcome, not causal effects, because breadth and contact diversity share papers.

Diffusion trajectories. No discrete trajectory typology passes the naming rule. DTW k -medoids and HMM-based clustering on state sequences give an adjusted Rand index of 0.222. PCA reveals a continuum: PC1 (38.8% of variance) is a breadth axis, PC2 (10.7%) is a keep-versus-lose axis. Early openness (OPEN) correlates with PC1 beyond B5 (held-out partial 0.120, $I^2 = 0$) but not with PC2 (partial -0.07 to -0.11).

Sequence and closure tests. A home-prominence half-peak vs off-home take-off test finds no signal beyond the mechanical lag (DEV excess -0.009 $[-0.015, -0.003]$, held-out $+0.011$ $[0.005, 0.016]$). Intersection-born concepts take off later (HR 0.47 $[0.42, 0.54]$). A within-concept closure test is null on DEV⁶ (density PPML $b = -0.070$ $[-0.180, 0.040]$; OPEN_{home} $b = +0.015$ $[-0.038, 0.069]$). On held-out fields, OPEN_{home} is opposite-signed (-0.079 $[-0.146, -0.013]$). Early openness is a between-concept trait, not a within-concept mechanism.

⁵Code: https://github.com/ai-inventor-papers/ai-invention-8795ea-concepts-spread-once-adopters-make-them/tree/fork/run_DVtwxCx0JbFq/round-4/experiment-12

⁶Code: https://github.com/ai-inventor-papers/ai-invention-8795ea-concepts-spread-once-adopters-make-them/tree/fork/run_DVtwxCx0JbFq/round-5/experiment-15




Figure 2: Retained-frontier coefficient d_0 (log-odds of entering a field per SD of relatedness to the off-home fields that currently retain the concept) from the conditional-logit entry model (rung R3, controlling for home relatedness, field size, entered density, own gateway centrality, RCA density D_{rca} and volume density D_{vol}), by unit. Circles with bars show point estimates with 95% concept-bootstrap confidence intervals. The light-blue row is the development (DEV) set used for selection. Dark-blue rows are the held-out domain groups and the 2010-14 onset cohort. The black diamond is the DerSimonian-Laird random-effects estimate pooled over the four held-out groups, and its width spans its 95% CI. The right column prints each estimate and interval; n is the number of concepts in each unit. The dashed line marks $d_0 = 0$. d_0 is positive with a CI excluding zero in all three evaluable held-out groups (Physical Sciences 0.148, Life & Env. 0.401, Social Sciences 0.297) and in the cohort (0.321). Math & Decision ($n = 165$, excluded as underpowered before the freeze) is null (0.065 [-0.110, 0.234]). The pooled estimate is 0.243 [0.118, 0.368] with high heterogeneity ($I^2 = 0.92$). The frozen verdict is nonetheless PARTIAL: the pre-registered volume-matched contrast is null, and the effect is specific to the PMI backbone.

3.2.4 Discussion of RQ2

The retained-frontier result extends the principle of relatedness to concept adoption. The positive coefficient is large and consistent across three of three evaluable groups, but the PARTIAL verdict limits its interpretation: we cannot fully separate persistence from sustained volume as a predictor of field entry. The effect is also backbone-specific: it holds on the sparse PMI backbone but not under the standard Hidalgo proximity. The breadth decomposition shows that integrating and localised concepts differ primarily in how many fields they contact early, not in how many more fields they enter later. The null closure test and opposite-signed held-out result indicate that openness does not operate as a within-concept mechanism driving diffusion.

3.3 Confound analysis

The openness signal could be an artefact of small yearly samples or of dynamic churn. We tested four null models⁷:

Fixed-size rarefaction. Subsampling to 10 home papers per year retains 68% of the pooled association. The bias-corrected Chao estimator retains 91%.

Within-concept permutation null. The excess churn over a 200-draw within-concept year-label permutation null is indistinguishable from zero: excess NOVCHURN pooled PSP = +0.008 [−0.02, +0.03], compared to raw NOVCHURN = +0.116 [+0.09, +0.14]. The R^2 of raw persistence on the permutation-null mean is 0.66. What predicts breadth is a static property of the concept’s home-topic partner mix.

Degree-preserving configuration null. Replacing raw ego density and persistence by curve-ball configuration z-scores raises OPEN_{home} from +0.092 to +0.115.

Reliability. Split-half reliability of NOVCHURN_{raw} is 0.48. The temporal-excess variants (V2) have split-half Spearman–Brown reliability of approximately 0.01–0.05, and the planted-churn check fails. The disattenuated pooled NOVCHURN PSP is 0.178 [0.141, 0.214] (approximate, CI not independently re-derived).

Because the V2 temporal-excess measure is unreliable, the confound analysis cannot adjudicate temporal churn at typical sample sizes (~10 home papers/year). The association that survives behaves like a static dispersion property. Temporal turnover is untested rather than refuted.

Partner decomposition. The home-only signal is carried by partners from new Leiden communities arriving through mixed-field papers: contrast $C_2 = +0.102$ [+0.069, +0.133] (Holm $p = 0.0025$), domain Shapley $\phi = +0.132$ on the cohort (method = +0.012). Controlling for the share of early bridging papers halves NOVCHURN from 0.118 to 0.056.

3.4 Exploratory AI-domain stage

An exploratory retrospective atlas of 37 AI and computer science concepts was constructed from the large panel, examining how concepts such as deep learning, reinforcement learning and generative adversarial networks evolved through the co-occurrence network before their cross-field diffusion. This atlas is outcome-selected and retrospective; it provides qualitative context but no hypothesis-testing power.

⁷Code: https://github.com/ai-inventor-papers/ai-invention-8795ea-concepts-spread-once-adopters-make-them/tree/fork/run_DVtwxCx0JbFq/round-5/experiment-16

3.5 Representative case studies

Seven matched case-study pairs were drawn from the breadth decomposition analysis (Table 5). These pairs are illustration only; all are descriptive with no p -value.

Table 5: Representative matched case-study pairs contrasting high-breadth and low-breadth concepts matched on early volume. All pairs are illustration; no inferential claim is made.

High-breadth concept	Low-breadth concept	Group
Graphics processing unit	Vertical-axis wind turbine	Eng
Shotgun proteomics	Image-guided radiation therapy	BGM
Nanocarriers	Nanosheet	CS
Soft power	Autonomous learning	SOC
Scopus	Oxygen reduction reaction	CS
Sclerostin	IgG 4-related disease	BGM
User-generated content	Mindfulness-based cognitive therapy	SOC

3.6 Dead ends and negative results

Table 6 lists the principal hypotheses and candidate signals that were tested and did not survive (across five iterations and 21 artifacts).

4 Discussion

4.1 What the openness index captures

The OPEN index captures the degree to which a concept’s early neighbourhood acquires partners from diverse communities rather than reinforcing existing ones. The signal is not demonstrably temporal turnover but a static topical-dispersion property. A within-concept permutation null absorbs the year-to-year component, yet the association persists: concepts whose home-field co-occurrence partners span many communities tend to reach more fields later. This is consistent with March’s (1991) exploration–exploitation distinction at the concept level. However, the within-concept closure test is null on DEV and opposite-signed on held-out fields, indicating that openness operates as a between-concept trait rather than as a within-concept causal mechanism.

The home-only signal ($\text{OPEN}_{\text{home}}$) is marginal. The stronger OPEN_{all} signal is partly mechanically coupled: early off-home papers that enter the ego network are a subset of the papers that determine the outcome. $\text{OPEN}_{\text{sizematch}}$ controls for this coupling and remains significant.

4.2 Reconciling consistency and breadth

Cheng et al. [2023] found that ideational consistency predicts a concept becoming “core.” We replicate their result for volume (+53%/SD) but show that, net of volume, the same property predicts *narrower* cross-field reach. The reversal holds in all five field groups ($I^2 = 0$) and replicates on a fresh cohort. The features predictive of local consolidation and growth are distinct from, and partly opposed to, the features predictive of broad disciplinary spread.




Figure 3: Illustrative matched case-study pair from Experiment 12 (case pair 1 of 7): graphics processing unit (GPU; onset $t_0 = 2008$, high early-neighbourhood openness OPEN) versus vertical-axis wind turbine (VAWT; $t_0 = 2009$, low OPEN). Both have an Engineering home field and are matched on early volume and growth. (a) Number of off-home venue fields each concept has ever entered (solid line, filled markers) and currently retains (dashed line, open markers) in each year from t_0 to $t_0 + 8$; GPU is shown in black and VAWT in grey on a shared axis. By $t_0 + 2$ GPU had entered 7 off-home fields and VAWT 2; by $t_0 + 8$ they had entered 11 and 6 and retained 8 and 4. (b, c) Topic co-occurrence ego networks over the first three years (t_0 to $t_0 + 2$, all papers). The nodes are the concept’s positive-PMI neighbour topics (the concept itself is not drawn), the edges are links between those topics in the 2010–14 topic backbone, node area scales with the number of the concept’s papers carrying the topic, and colour gives the topic’s Leiden community, named by the plurality OpenAlex field of its topics (plus the runner-up field when it has at least 0.6 times as many). GPU’s neighbourhood has 34 topics in 8 communities (density 0.30), led by Engineering & Physics, Computer Science and Medicine communities. VAWT’s has 12 topics in a single Engineering community (density 0.70). The pair illustrates the openness–breadth association and is not evidence for it (case_pairs.json, pair 1).

Table 6: Principal dead ends and negative results across five research iterations.

Hypothesis / candidate	Decisive test	Result
Naturalisation gap (A_h^*)	Exp. 1	Fails every clause (-0.006)
$D_{\text{ratio}}, F_{\text{res}}$	Exp. 3	No candidate survives
Gateway landing on O_{2r}	Exp. 4, 5	$\Delta\text{AUC } -0.00001$ held-out
H3 gateway-weighted landing	Exp. 5	Pooled CI includes 0 $[-0.006, 0.065]$
Ordering (central-first)	Exp. 6, Eval. 2	MIXED: 32.6% of broad concepts; negative lead-lag; pre-trend
Two discrete trajectory classes	Exp. 6, 12	HMM-vs-DTW ARI 0.094/0.222
Abandonment penalty (d_{lost})	Exp. 7	Specification-dependent: A1 -0.007 , R4 $+0.064$, min-cp -0.030
Volume-matched retained-frontier contrast	Exp. 7	Null: -0.028 $[-0.105, +0.046]$ held-out
O_5 external recognition	Exp. 8, Eval. 2	No indicator beats B5 + onset year
Rescue and relay models	Exp. 6	Not supported
Within-concept closure	Exp. 11, 15C	Null on DEV; opposite-signed held-out (-0.079)
$n_{\text{authors_early}}$ for O_3	Exp. 8, 10	Non-replication on fresh cohort ($+0.014$)
RETENTION_RATIO at R2/R3	Exp. 10	Null on cohort (-0.043 $[-0.116, 0.031]$)
Trait prediction (P-B1)	Exp. 15B	Fails: yearly OPEN _{home} ICC 0.34 – 0.39
Frame N Cheng reversal	Exp. 13	Not confirmed (-0.064 $[-0.159, 0.031]$)
Co-author candidate S	Exp. 8	Computed; not confirmed for any outcome

4.3 The retained-frontier claim

The confirmed field-entry result ($d_0 = 0.322$ on the independent frame) extends the principle of relatedness from economic geography [Hidalgo et al., 2007, Neffke et al., 2011] to concept adoption. The result is PARTIAL: the volume-matched contrast is null, and the effect is backbone-specific (it reverses under minimum-conditional-probability proximity). What survives is a PMI-backbone relative-odds effect that cannot be separated from volume. The I^2 of 0.92 indicates substantial cross-domain heterogeneity, with Life & Environment showing the strongest signal and Mathematics & Decision Sciences null.

4.4 Limitations

Effect size. The OPEN–breadth association is modest. $\text{OPEN}_{\text{home}}$ adds no forecasting gain over B5 (+0.002 [−0.003, +0.008]).

Measurement noise. Split-half reliability of $\text{OPEN}_{\text{home}}$ is 0.49. The temporal-excess variants have split-half reliability below 0.05, so temporal versus static dispersion cannot be resolved at available sample sizes.

Vocabulary bias. The legacy OpenAlex concept vocabulary is survivor-selected. Frame N is small (636 concepts) with pre-unseal power of 0.47.

Coupling. About half of the OPEN_{all} signal is mechanical coupling between the ego-network construction (which uses off-home papers) and the breadth outcome. $\text{OPEN}_{\text{home}}$ is free of this coupling but marginal.

Domain heterogeneity. Life & Environment shows the weakest OPEN signal (PSP +0.042). The retained-field relatedness I^2 is 0.92.

Not causal. Early ego-network openness could reflect a concept’s intrinsic generality, the diversity of the research community around it, or the breadth of the problems it addresses.

Prior unsealing. Held-out outcomes were previously unsealed by two earlier experiments in the same project, so the indicator-screen held-out results are robustness checks.

5 Conclusions

We screened 53 early temporal network indicators for predicting the cross-disciplinary breadth of new scientific concepts, using 12,499 concepts from the OpenAlex corpus with held-out validation and a fresh cohort confirmation. Seven indicators are confirmed, composing into an openness index that captures topical non-redundancy in a concept’s early co-occurrence neighbourhood. The signal is a static between-concept trait; temporal turnover could not be tested at available sample sizes. Early openness predicts breadth; early consolidation predicts volume but not breadth. Concepts spread next to fields related to those currently retaining them, but the pre-registered verdict is PARTIAL: persistence and volume are confounded, and the effect is backbone-specific.

Declarations

Funding

Not applicable.

Conflict of interest

The authors declare no conflict of interest.

Availability of data and materials

The study uses the publicly available OpenAlex bulk snapshot (2026-09-23). Concept frames, indicator matrices and analysis scripts are available from the corresponding author upon request.

Code availability

Analysis code is available from the corresponding author upon request.

Authors’ contributions

All authors contributed to the study design, analysis and manuscript preparation.

References

- R. Burt. Structural holes and good ideas¹. In *American Journal of Sociology*, volume 110, pages 349 – 399, 2004. doi: 10.1086/421787.
- M. Callon, J. Courtial, and F. Laville. Co-word analysis as a tool for describing the network of interactions between basic and technological research: The case of polymer chemistry. In *Scientometrics*, volume 22, pages 155–205, 1991. doi: 10.1007/BF02019280.
- Damon Centola. The spread of behavior in an online social network experiment. In *Science*, volume 329, pages 1194 – 1197, 2010. doi: 10.1126/science.1185231.
- David Chavalarias and Jean-Philippe Cointet. Phylomemetic patterns in science evolution—the rise and fall of scientific fields. In *PLoS ONE*, volume 8, 2013. doi: 10.1371/journal.pone.0054847.
- M. Cheng, D. Smith, Xiang Ren, Hancheng Cao, Sanne Smith, and Daniel A. McFarland. How new ideas diffuse in science. In *American Sociological Review*, volume 88, pages 522 – 561, 2023. doi: 10.1177/00031224231166955.
- H. Cherifi, G. Palla, B. Szymanski, and Xiaoyan Lu. On community structure in complex networks: challenges and opportunities. In *Applied Network Science*, volume 4, 2019. doi: 10.1007/s41109-019-0238-9.
- R. Dersimonian and N. Laird. Meta-analysis in clinical trials. In *Controlled Clinical Trials*, volume 7 3, pages 177–88, 1986. doi: 10.1016/0197-2456(86)90046-2.
- M. De Domenico, E. Omodei, and A. Arenas. Quantifying the diaspora of knowledge in the last century. In *Applied Network Science*, volume 1, 2016. doi: 10.1007/s41109-016-0017-9.
- J. Fitzgerald, Sanna Ojanperä, and N. O’Clery. Is academia becoming more localised? the growth of regional knowledge networks within international research collaboration. In *Applied Network Science*, volume 6, 2021. doi: 10.1007/s41109-021-00371-w.
- S. Fortunato, Carl T. Bergstrom, K. Börner, James A. Evans, D. Helbing, et al. Science of science. In *Science*, volume 359, 2018. doi: 10.1126/science.aao0185.
- Yuan Gao, Zhen-Yan Zhu, Raja Kali, and M. Riccaboni. Community evolution in patent networks: technological change and network dynamics. In *Applied Network Science*, volume 3, 2018. doi: 10.1007/s41109-018-0090-3.

- César A. Hidalgo, B. Klinger, A. Barabási, and R. Hausmann. The product space conditions the development of nations. In *Science*, volume 317, pages 482 – 487, 2007. doi: 10.1126/science.1144581.
- Julian P. T. Higgins and S. Thompson. Quantifying heterogeneity in a meta-analysis. In *Statistics in Medicine*, volume 21, 2002. doi: 10.1002/SIM.1186.
- Bas Hofstra, V. V. Kulkarni, Sebastian Munoz-Najar Galvez, Bryan He, Dan Jurafsky, and Daniel A. McFarland. The diversity–innovation paradox in science. In *Proceedings of the National Academy of Sciences of the United States of America*, volume 117, pages 9284 – 9291, 2019. doi: 10.1073/pnas.1915378117.
- J. March. Exploration and exploitation in organizational learning. In *STUDI ORGANIZZATIVI*, 1991. doi: 10.1287/ORSC.2.1.71.
- Frank M. H. Neffke, M. Henning, and R. Boschma. How do regions diversify over time? industry relatedness and the development of new growth paths in regions. *Economic Geography*, 87:237 – 265, 2011. doi: 10.1111/j.1944-8287.2011.01121.x.
- Joseph D. O’Brien, I. Dassios, and J. Gleeson. Spreading of memes on multiplex networks. In *New Journal of Physics*, volume 21, 2018. doi: 10.1088/1367-2630/ab05ef.
- B. Renoust, Vivek Claver, and Jean-François Baffier. Multiplex flows in citation networks. In *Applied Network Science*, volume 2, 2017. doi: 10.1007/s41109-017-0035-2.
- D. Rotolo, D. Hicks, and Ben Martin. What is an emerging technology? *arXiv preprint arXiv:1503.00673*, 2015. doi: 10.1016/J.RESPOL.2015.06.006.
- Angelo Salatino, Francesco Osborne, and E. Motta. How are topics born? understanding the research dynamics preceding the emergence of new areas. In *PeerJ Computer Science*, volume 3, page e119, 2017. doi: 10.7287/peerj.preprints.2306v1.
- N. Sattar, Aydın Buluç, Khaled Z. Ibrahim, and S. Arifuzzaman. Exploring temporal community evolution: algorithmic approaches and parallel optimization for dynamic community detection. In *Applied Network Science*, volume 8, pages 1–29, 2023. doi: 10.1007/s41109-023-00592-1.
- J. Ugander, L. Backstrom, Cameron A. Marlow, and J. Kleinberg. Structural diversity in social contagion. In *Proceedings of the National Academy of Sciences of the United States of America*, volume 109, pages 5962 – 5966, 2012. doi: 10.1073/pnas.1116502109.
- L. Weng, F. Menczer, and Yong-Yeol Ahn. Virality prediction and community structure in social networks. In *Scientific Reports*, volume 3, 2013. doi: 10.1038/srep02522.